

CS122A: Introduction to Data Management

Lecture #7 *Relational Algebra I*

Instructor: Chen Li

Relational Query Languages

- ❖ Query languages: Allow manipulation and **retrieval of data** from a database.
- ❖ Relational model supports simple, powerful QLs:
 - Strong formal foundation based on logic.
 - Allows for much optimization.
- ❖ Query Languages **≠** programming languages!
 - QLs not expected to be “Turing complete”.
 - QLs not intended to be used for complex calculations.
 - QLs support easy, efficient access to large data sets.

Formal Relational Query Languages

- ❖ Two mathematical Query Languages form the basis for “real” languages (e.g., SQL), and for implementation:
 - Relational Algebra: More **operational**, very useful for representing execution plans.
 - Relational Calculus: Let's users describe what they want, rather than how to compute it.
(**Non-operational**, declarative.)

Preliminaries

- ❖ A query is applied to *relation instances*, and the result of a query is also a relation instance.
 - *Schemas of input* relations for a query are *fixed* (but query will run regardless of instance!)
 - The *schema for the result* of a given query is also *fixed*! Determined by definition of query language constructs.
- ❖ Positional vs. named-field notation:
 - Positional notation easier for formal definitions, named-field notation more readable.
 - Both used in SQL (but try to avoid positional stuff!)

Example Instances

R1

| <u>sid</u> | <u>bid</u> | <u>day</u> |
|------------|------------|------------|
| 22 | 101 | 10/10/96 |
| 58 | 103 | 11/12/96 |

- ❖ “Sailors” and “Reserves” relations for our examples.
- ❖ We’ll use positional or named field notation, and assume that names of fields in query results are “inherited” from names of fields in query input relations.

S1

| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

S2

| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 28 | yuppy | 9 | 35.0 |
| 31 | lubber | 8 | 55.5 |
| 44 | guppy | 5 | 35.0 |
| 58 | rusty | 10 | 35.0 |

Relational Algebra

❖ Basic operations:

- Selection (σ) Selects a subset of rows from relation.
- Projection (π) Deletes unwanted columns from relation.
- Cross-product (\times) Allows us to combine two relations.
- Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
- Union (\cup) Tuples in reln. 1 and in reln. 2.

❖ Additional operations:

- Intersection, join, division, renaming: Not essential, but (very!) useful. (I.e., don't add expressive power, but...)

❖ Since each operation returns a relation, **operations can be composed!** (Algebra is “closed”.)

Projection

- ❖ Removes attributes that are not in *projection list*.
- ❖ *Schema* of result contains exactly the fields in the projection list, with the same names that they had in the (only) input relation.
- ❖ Relational projection operator has to eliminate *duplicates*! (Why??)
 - Note: real systems typically don't do duplicate elimination unless the user explicitly asks for it. (Q: Why not?)

| sname | rating |
|--------|--------|
| yuppy | 9 |
| lubber | 8 |
| guppy | 5 |
| rusty | 10 |

$\pi_{sname, rating}(S2)$

| age |
|------|
| 35.0 |
| 55.5 |

$\pi_{age}(S2)$

Selection

| sid | sname | rating | age |
|-----|-------|--------|------|
| 28 | yuppy | 9 | 35.0 |
| 58 | rusty | 10 | 35.0 |

- ❖ Selects rows that satisfy a *selection condition*.
- ❖ No duplicates in result! (Why?)
- ❖ *Schema* of result identical to schema of its (only) input relation.
- ❖ *Result* relation can be the *input* for another relational algebra operation! (This is *operator composition*.)

$$\sigma_{rating > 8}(S2)$$

| sname | rating |
|-------|--------|
| yuppy | 9 |
| rusty | 10 |

$$\pi_{sname, rating}(\sigma_{rating > 8}(S2))$$

Union, Intersection, Set-Difference

- ❖ All of these operations take two input relations, which must be union-compatible:
 - Same number of fields.
 - “Corresponding” fields are of the same type.
- ❖ What is the *schema* of result?

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | dustin | 7 | 45.0 |

$S1 - S2$

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |
| 44 | guppy | 5 | 35.0 |
| 28 | yuppy | 9 | 35.0 |

$S1 \cup S2$

| sid | sname | rating | age |
|-----|--------|--------|------|
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

$S1 \cap S2$

Q: Any issues w/ duplicates?

Cross-Product

- ❖ Each row of S1 is paired with each row of R1.
- ❖ *Result schema* has one field per field of S1 and R1, with field names 'inherited' if possible.
 - *Conflict*: Both S1 and R1 have a field called *sid*.

| (sid) | sname | rating | age | (sid) | bid | day |
|-------|--------|--------|------|-------|-----|----------|
| 22 | dustin | 7 | 45.0 | 22 | 101 | 10/10/96 |
| 22 | dustin | 7 | 45.0 | 58 | 103 | 11/12/96 |
| 31 | lubber | 8 | 55.5 | 22 | 101 | 10/10/96 |
| 31 | lubber | 8 | 55.5 | 58 | 103 | 11/12/96 |
| 58 | rusty | 10 | 35.0 | 22 | 101 | 10/10/96 |
| 58 | rusty | 10 | 35.0 | 58 | 103 | 11/12/96 |

Result
relation name

Attribute
renaming list

Source
expression
(anything!)

- Renaming operator: $\rho (C(1 \rightarrow sid1, 5 \rightarrow sid2), S1 \times R1)$

Renaming

❖ *Conflict*: S1 and R1 both had *sid* fields, giving:

| (sid) | sname | rating | age | (sid) | bid | day |
|-------|--------|--------|------|-------|-----|----------|
| 22 | dustin | 7 | 45.0 | 22 | 101 | 10/10/96 |
| ... | ... | ... | ... | ... | ... | ... |
| 58 | rusty | 10 | 35.0 | 58 | 103 | 11/12/96 |

❖ Several renaming options available:

$\rho (S1R1(1 \rightarrow sid1), S1 \times R1)$ ← Positional renaming

$\rho (TempS1(sid \rightarrow sid1), S1)$ ← Name-based renaming

$TempS1 \times R1$

$(\pi_{sid \rightarrow sid1, sname, rating, age}(S1)) \times R1$ ← Generalized projection
(I like this notation best ☺)

Joins

❖ Condition Join: $R \bowtie_c S = \sigma_c (R \times S)$

| (sid) | sname | rating | age | (sid) | bid | day |
|-------|--------|--------|------|-------|-----|----------|
| 22 | dustin | 7 | 45.0 | 58 | 103 | 11/12/96 |
| 31 | lubber | 8 | 55.5 | 58 | 103 | 11/12/96 |

$$S1 \bowtie_{sid < sid} R1$$

- ❖ *Result schema* same as that of cross-product.
- ❖ Fewer tuples than cross-product, so might be able to compute more efficiently
- ❖ Sometimes (often!) called a *theta-join*.

More Joins

- ❖ Equi-Join: A special case of condition join where the condition c contains only *equalities*.

| sid | sname | rating | age | bid | day |
|-----|--------|--------|------|-----|----------|
| 22 | dustin | 7 | 45.0 | 101 | 10/10/96 |
| 58 | rusty | 10 | 35.0 | 103 | 11/12/96 |

$$S1 \bowtie_{sid} R1$$

- ❖ Result schema similar to cross-product, but only one copy of fields for which equality is specified.
- ❖ Natural Join: An equijoin on *all* commonly named fields.

Division

- ❖ Not a primitive operator, but extremely useful for expressing queries like:
Find sailors who have reserved all boats.
- ❖ Let A have 2 fields, x and y , while B has one field y , so we have relations $A(x,y)$ and $B(y)$:
 - **A/B contains the x tuples (e.g., sailors) such that for every y tuple (e.g., boat) in B , there is an xy tuple in A .**
 - Or: If the set of y values (boats) associated with an x value (sailor) in A contains all y values in B , the x value is in A/B .
- ❖ In general, x and y can be any lists of fields; y is the list of fields in B , and $x \cup y$ is the list of fields of A .

Examples of Division A/B

| sno | pno |
|-----|-----|
| s1 | p1 |
| s1 | p2 |
| s1 | p3 |
| s1 | p4 |
| s2 | p1 |
| s2 | p2 |
| s3 | p2 |
| s4 | p2 |
| s4 | p4 |

A

| pno |
|-----|
| p2 |

B1



| sno |
|-----|
| s1 |
| s2 |
| s3 |
| s4 |

A/B1

| pno |
|-----|
| p2 |
| p4 |

B2



| sno |
|-----|
| s1 |
| s4 |

A/B2

| pno |
|-----|
| p1 |
| p2 |
| p4 |

B3



| sno |
|-----|
| s1 |

A/B3

Expressing A/B Using Basic Operators

(Advanced Topic ☺)

- ❖ Division not an essential op; just a useful shorthand.
(Also true of joins, but joins are so common and important that relational database systems implement joins specially.)
- ❖ *Idea*: For A/B , compute all x values that are not “disqualified” by some y value in B .
 - x value is *disqualified* if by attaching a y value from B , we obtain an xy tuple that does not appear in A .

Disqualified x values (D): $\pi_x ((\pi_x(A) \times B) - A)$

A/B : $\pi_x(A) - D$

Ex: Wisconsin Sailing Club Database

Sailors

| sid | sname | rating | age |
|-----|---------|--------|------|
| 22 | Dustin | 7 | 45.0 |
| 29 | Brutus | 1 | 33.0 |
| 31 | Lubber | 8 | 55.5 |
| 32 | Andy | 8 | 25.5 |
| 58 | Rusty | 10 | 35.0 |
| 64 | Horatio | 7 | 35.0 |
| 71 | Zorba | 10 | 16.0 |
| 74 | Horatio | 9 | 35.0 |
| 85 | Art | 4 | 25.5 |
| 95 | Bob | 3 | 63.5 |

Reserves

| sid | bid | date |
|-----|-----|----------|
| 22 | 101 | 10/10/98 |
| 22 | 102 | 10/10/98 |
| 22 | 103 | 10/8/98 |
| 22 | 104 | 10/7/98 |
| 31 | 102 | 11/10/98 |
| 31 | 103 | 11/6/98 |
| 31 | 104 | 11/12/98 |
| 64 | 101 | 9/5/98 |
| 64 | 102 | 9/8/98 |
| 74 | 103 | 9/8/93 |

Boats

| bid | bname | color |
|-----|-----------|-------|
| 101 | Interlake | blue |
| 102 | Interlake | red |
| 103 | Clipper | green |
| 104 | Marine | red |

Find names of sailors who've reserved boat #103

Sailors(sid, sname, rating, age)
Boats(bid, bname, color)

Reserves(sid, bid, day)

❖ Solution 1: $\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$

❖ Solution 2: $\rho(Temp1, \sigma_{bid=103} Reserves)$

$\rho(Temp2, Temp1 \bowtie Sailors)$

$\pi_{sname}(Temp2)$

❖ Solution 3: $\pi_{sname}(\sigma_{bid=103}(Reserves \bowtie Sailors))$

Ex: Wisconsin Sailing Club Database

$\sigma_{bid=103}$ Reserves

| sid | bid | date |
|-----|-----|---------|
| 22 | 103 | 10/8/98 |
| 31 | 103 | 11/6/98 |
| 74 | 103 | 9/8/93 |

$\pi_{sname}((\sigma_{bid=103} \text{Reserves}) \bowtie \text{Sailors})$

$(\sigma_{bid=103} \text{Reserves}) \bowtie \text{Sailors}$

| sid | bid | date | sname | rating | age |
|-----|-----|---------|---------|--------|------|
| 22 | 103 | 10/8/98 | Dustin | 7 | 45.0 |
| 31 | 103 | 11/6/98 | Lubber | 8 | 55.5 |
| 74 | 103 | 9/8/93 | Horatio | 9 | 35.0 |

| sname |
|---------|
| Dustin |
| Lubber |
| Horatio |

Find names of sailors who've reserved a red boat

Sailors(sid, sname, rating, age)

Reserves(sid, bid, day)

Boats(bid, bname, color)

- ❖ Information about boat color only available in Boats; so need to do a join:

$$\pi_{sname}((\sigma_{color='red'} Boats) \bowtie Reserves \bowtie Sailors)$$

- ❖ A more “efficient” solution:

$$\pi_{sname}(\pi_{sid}((\pi_{bid} \sigma_{color='red'} Boats) \bowtie Res) \bowtie Sailors)$$

A query optimizer can find the latter given the first solution!

Find sailors who've reserved a red or a green boat

Sailors(sid, sname, rating, age)

Reserves(sid, bid, day)

Boats(bid, bname, color)

- ❖ Can identify all red or green boats, then find sailors who've reserved one of these boats:

$$\rho \text{ (Tempboats, } (\sigma_{color='red' \vee color='green'} \text{Boats}))$$
$$\pi_{sname}(\text{Tempboats} \bowtie \text{Reserves} \bowtie \text{Sailors})$$

- ❖ Can also define Tempboats using union! (Q: How?)
- ❖ What happens if \vee is replaced by \wedge in this query?

Find sailors who've reserved a red and a green boat

Sailors(sid, sname, rating, age)

Reserves(sid, bid, day)

Boats(bid, bname, color)

- ❖ Previous approach won't work! Must identify sailors who've reserved red boats and sailors who've reserved green boats, then find their intersection (note that *sid* is a key for Sailors):

$$\rho (Tempred, \pi_{sid}((\sigma_{color='red'} Boats) \bowtie Reserves))$$
$$\rho (Tempgreen, \pi_{sid}((\sigma_{color='green'} Boats) \bowtie Reserves))$$
$$\pi_{sname}((Tempred \cap Tempgreen) \bowtie Sailors)$$

Find the names of sailors who've reserved all boats

Sailors(sid, sname, rating, age)

Reserves(sid, bid, day)

Boats(bid, bname, color)

- ❖ Uses division; schemas of the input relations feeding / operator must be *carefully chosen*:

$$\rho \text{ (Tempsids, } (\pi_{sid,bid} \text{Reserves}) / (\pi_{bid} \text{Boats}))$$

$$\pi_{sname} (\text{Tempsids} \bowtie \text{Sailors})$$

- ❖ To find sailors who've reserved all 'Interlake' boats:

$$\dots / \pi_{bid} (\sigma_{bname='Interlake'} \text{Boats})$$

Relational Algebra Summary

- ❖ The relational model has (several) rigorously defined query languages that are both simple and powerful in nature.
- ❖ Relational algebra is more operational; very useful as an internal representation for query evaluation plans.
- ❖ Several ways of expressing a given query; a query optimizer should choose the most efficient version. (Take CS122C? ☺)
- ❖ We'll add a few more operators later on...
- ❖ Next up for now: *Relational Calculus*